

Threat Investigation Report

Identity: u411@a.alpineskihouse.co (Denise Charron)

Microsoft Sentinel — Exposure Graph & Behavioral Analysis

Classification: CONFIDENTIAL — For Customer Distribution
Prepared by: Security Operations Center
Date: March 4, 2026
Workspace: lawsentinel1 — Alpine Ski House tenant
Incident Reference: Insider Risk Incident #49130

Table of Contents

- 1. [Executive Summary](#)
- 2. [Identity Profile](#)
- 3. [Sign-in Behavioral Analysis](#)
- 4. [Exposure Graph — Blast Radius](#)
- 5. [Hunting Over Graph — Key Findings](#)
- 6. [Critical Resource Exposure Detail](#)
- 7. [Risk Assessment Matrix](#)
- 8. [Immediate Recommended Actions](#)

Executive Summary

During a threat hunting engagement on the Alpine Ski House Microsoft Sentinel workspace, the identity u411@a.alpineskihouse.co (**Denise Charron**) was identified as a **critical insider risk subject** requiring urgent investigation.

Leveraging the **Microsoft Sentinel Exposure Graph**, we mapped every Azure resource reachable from this identity. The results reveal an extraordinarily broad blast radius spanning:

- **17 Key Vaults** (Criticality: HIGH) — including vaults holding SOC secrets, network security credentials, and threat intelligence key material
- **30+ Storage Accounts** — including `midnightblizzardstorage`, a resource under a resource group named `midnightblizzard_infra`, strongly associated with Russian state-sponsored threat actor **Midnight Blizzard (NOBELIUM)**
- **8 Kubernetes (AKS) Clusters** — full administrative surface of the container environment
- **12 AI/Cognitive Services** — including 3 Azure OpenAI instances and custom agent systems in the `zavabot` resource group
- **6 Machine Learning Workspaces** — including `CyberSOC_ML`, providing access to the SOC's own detection models

This account additionally shares infrastructure linkage with a second compromised identity (`elviala@zava-corp.com`) via the `zavabot` AI agent resource group — suggesting a **coordinated insider threat or supply chain compromise scenario**.

 **This account is NOT enrolled in Microsoft Entra ID Identity Protection.** No risk signals are being generated in `AADRiskyUsers` , creating a critical detection blind spot.

Identity Profile

Field	Value
User Principal Name	<code>u411@a.alpineskihouse.co</code>
Display Name	Denise Charron
AAD Object ID	<code>c4e33ecf-bbb2-4948-953d-d3bfe73a7d74</code>
Exposure Graph Node ID	<code>0e447a7d9dcd433797d19b21a0d75663</code>
Criticality Tag	Critical Asset (tagged in Exposure Graph)
Identity Protection Status	 Not enrolled — blind spot
Associated Incident	Insider Risk Incident #49130
Tenant	Alpine Ski House (<code>alpineskihouse.co</code>)
Investigation Date	March 4, 2026

Sign-in Behavioral Analysis

Data source: `SignInLogs` — 30-day lookback (Feb 2 – Mar 4, 2026)

Summary Statistics

Metric	Value	Significance
Total Sign-ins	800	High-frequency access
Failed Sign-ins	278 (34.75%)	 Elevated — credential stuffing, MFA fatigue, or automation
Unique Source IPs	6	Concentrated — low variety suggests scripted/automated access
Geographic Locations	United States only	No anomalous geo-hop
Unique Applications	35	 Very high — includes admin, API, and CLI tools

First Seen	Feb 2, 2026	
Last Seen	Mar 4, 2026	Active at time of investigation

Source IP Addresses

20.97.10.99

– Microsoft Azure Service IP

52.167.112.163

– Microsoft Azure

52.177.6.198

– Microsoft Azure

52.167.115.14

– Microsoft Azure

52.177.6.190

– Microsoft Azure

76.182.132.142

– Residential/ISP IP ← NON-CORPORATE

⚠ IP 76.182.132.142 is a non-corporate residential ISP address. Sign-ins from this IP should be reviewed for context.

High-Risk Application Activity


The following privileged applications were used during the investigation window — consistent with **manual or scripted administrative reconnaissance and data access**:


Application	Risk Indicator
Azure Portal	Full UI administrative access
Azure AI Studio App	Direct AI workload access
Microsoft Azure PowerShell	● Scripted automation capability
Microsoft Azure CLI	● Scripted automation capability
Microsoft Graph Command Line Tools	● Full tenant API programmatic access
Graph Explorer	Directory and resource enumeration
Microsoft 365 Security and Compliance Center	Access to compliance/DLP/audit logs
Teams Admin Portal Service	Tenant-wide Teams administration
Microsoft 365 Admin portal	Global tenant admin portal


Analysis: The combination of CLI tools, PowerShell, and Graph API access alongside a 34.75% failure rate is consistent with either automated tooling or a compromised credential being actively used by a threat actor alongside the legitimate user.


Exposure Graph — Blast Radius

Methodology: The Microsoft Sentinel Exposure Graph was queried using `graph_find_blastradius("Denise Charron")`. This tool traverses all walkable edges from the identity node and returns every Azure resource reachable via direct or indirect permission paths. All paths identified were **1-hop (direct permissions)** — meaning no privilege escalation is required.









Reading the Graph

The diagram above represents the **blast radius** of the identity `u411@a.alpineskihouse.co` as computed by the Microsoft Sentinel Exposure Graph. Each arrow labeled **"direct permissions"** means:

"If this account is compromised, the attacker can access this resource with zero additional steps — no privilege escalation, no lateral movement required."

Every resource in the graph is a one-hop path from the user node. This is significant: for most identities, a large blast radius requires multiple hops (e.g., user → service principal → managed identity → resource). Here, **all 60+ resources are directly reachable in a single step**, meaning the account has been granted broad Azure RBAC roles directly.

Color/shape coding:

-  **Key Vaults (purple/red)** — Highest criticality. Compromise gives attacker all stored secrets, API keys, connection strings, and certificates
-  **AKS Clusters (blue)** — Container orchestration clusters; compromise enables container escape, data exfiltration, or crypto-mining workloads
-  **AI Services (teal)** — Azure OpenAI and custom agent infrastructure; compromise enables prompt injection, model poisoning, or unauthorized LLM usage
-  **ML Workspaces (green)** — Training data, models, and pipelines; `CyberSOC_ML` contains the SOC's own threat detection models — a particularly sensitive target
-  **Storage Accounts (orange)** — Data at rest including corporate data, SAP exports, M365 audit logs, and the anomalous `midnightblizzardstorage`

Hunting Over Graph — Key Findings

The following findings emerged from cross-referencing the Exposure Graph blast radius with resource metadata, threat intelligence, and behavioral signals.

● Finding #1 — `midnightblizzardstorage` in `midnightblizzard_infra` Resource Group

Severity: CRITICAL

The storage account `midnightblizzardstorage` resides in a resource group named `midnightblizzard_infra`. **Midnight Blizzard (also known as NOBELIUM, APT29, Cozy Bear)** is a Russian SVR-affiliated advanced persistent threat actor responsible for the SolarWinds supply chain attack and multiple Microsoft/government breaches.

The presence of a resource group named after this threat actor — and its direct accessibility by an insider risk subject — warrants immediate escalation. Possible interpretations:

- The resource group was intentionally named by the threat actor as part of a controlled infrastructure deployment inside the tenant
- The naming is a deliberate honeypot or simulation environment (common in SOC demo tenants)
- The user has explicit permissions to threat actor-attributed infrastructure

Regardless of origin, an insider risk subject having direct access to infrastructure named `midnightblizzard_infra` is an unacceptable risk posture and must be investigated immediately.

● Finding #2 — Key Vault `MDTIACTOR` (RG: `cybersoc`)

Severity: CRITICAL

The Key Vault named `MDTIACTOR` is stored in the `cybersoc` resource group. The name `MDTIACTOR` strongly suggests this vault contains key material related to **Threat Actor simulation or production MDI/MDTI (Microsoft Defender Threat Intelligence) actor configurations**. Direct read access to this vault by `u411@a.alpineskihouse.co` could expose:

- Threat actor infrastructure credentials used in red team exercises
 - Production MDTI API keys or actor investigation data
 - Keys used to authenticate threat actor emulation tooling
-

● Finding #3 — SOC Infrastructure Key Vaults (SentinelDemo , cybersocm17693298145)

Severity: HIGH

Denise Charron has direct permissions to two Key Vaults in the cybersoc resource group:

- SentinelDemo — likely contains API keys, workspace tokens, or automation credentials for the Sentinel environment
- cybersocm17693298145 — associated with the SOC's machine learning pipeline

An adversary who controls SentinelDemo could potentially:

- Extract Sentinel workspace keys
 - Disable or modify analytics rules
 - Read or tamper with SIEM configuration
-

● Finding #4 — CyberSOC_ML Machine Learning Workspace

Severity: HIGH

The ML workspace CyberSOC_ML contains the SOC's own machine learning-based detection models. If an adversary accesses this workspace, they can:

- **Read** existing detection models — understanding exactly what behaviors trigger alerts, enabling evasion
- **Modify** or poison training data — causing the SOC to generate false positives or miss true positives
- **Delete** trained models — degrading detection capability

This represents a meta-attack on the defender's detection capability itself.

● Finding #5 — sharedcredsash and shellalpinesa (Cloud Shell Storage)

Severity: HIGH

- sharedcredsash Key Vault likely stores shared team credentials, service account passwords, or automation secrets
- shellalpinesa is Azure Cloud Shell storage, which persists PowerShell scripts, .profile configurations, SSH keys, and cached Azure CLI tokens between sessions

An adversary with access to Cloud Shell storage can recover cached credentials without triggering additional authentication events, making this a **credential harvesting target**.

● Finding #6 — zavabot Resource Group (AI Agent Infrastructure)

Severity: HIGH — Cross-Identity Linkage

Denise Charron (u411@a.alpineskihouse.co) has direct permissions to three resources in the zavabot resource group:

- zavainventoryplanning
- customdesigneragent
- zavacore-app

This is the same `zavabot` infrastructure flagged in the parallel investigation of `elviala@zava-corp.com` (Zava Corp insider risk). The cross-tenant linkage between Alpine Ski House user `u411` and Zava Corp AI agents suggests:

- Shared service account or managed identity granting cross-tenant access
- A common supply chain or integration point across both organizations
- Coordinated insider activity by individuals with access spanning both tenants

This cross-identity graph intersection is the most significant structural finding of this investigation and warrants a combined incident response across both organizations.

Finding #7 — Azure OpenAI Over-Permission (3 Instances)

Severity: MEDIUM

Three separate Azure OpenAI Cognitive Service instances are directly accessible:

- `mdc-d4ai-aoai-1`
- `cog-3wz76o5tfmbes`
- `cog-2o4jt2d674qo4`

A compromised identity with OpenAI access could send unauthorized requests to production LLMs, extract fine-tuned model configurations, or perform prompt injection attacks against downstream applications consuming these endpoints.

Finding #8 — Identity Protection Blind Spot

Severity: MEDIUM — Detection Gap

`AADRiskyUsers` returned **zero records** for `u411@a.alpineskihouse.co`. This means:

- No risk score is being generated for this identity
- No Conditional Access policies based on risk level will trigger
- No automatic remediation (force password reset, block sign-in) is occurring
- Standard SIEM detections that rely on `AADRiskyUsers` will not fire for this account


This is a significant gap in the detection coverage for what is otherwise a critical insider risk subject.

Critical Resource Exposure Detail

Resource	Type	Resource Group	Criticality	Why It Matters
<code>midnightblizzardstorage</code>	Storage	<code>midnightblizzard_infra</code>	CRITICAL	Threat actor-named infrastructure
<code>MDTIACTOR</code>	Key Vault	<code>cybersoc</code>	CRITICAL	Threat actor key material

SentinelDemo	Key Vault	cybersoc	CRITICAL	SIEM workspace secrets
sharedcredsash	Key Vault	democonfig	HIGH	Shared credential store
kv-netsec-alpineskihouse	Key Vault	N/A	HIGH	Network security credentials
CyberSOC_ML	ML Workspace	N/A	HIGH	SOC detection models — meta-attack surface
cybersocml17693298145	Key Vault	cybersoc	HIGH	SOC ML pipeline credentials
shellalpinesa	Storage	N/A	HIGH	Cloud Shell — cached tokens & scripts
zavainventoryplanning	Cognitive	zavabot	HIGH	Cross-tenant linkage to elviaa case
customdesigneragent	Cognitive	zavabot	HIGH	Cross-tenant linkage to elviaa case
zavacore-app	Cognitive	zavabot	HIGH	Cross-tenant linkage to elviaa case
sentitivestorage	Storage	xspm_demoresources	HIGH	Named sensitive storage
NetSec-kv-h09g	Key Vault	N/A	HIGH	Network security keys
NetSec-Demo-kv-pmo8	Key Vault	N/A	HIGH	Network security keys
alpine-mdc-vault	Key Vault	mdc	MEDIUM	Core MDC infrastructure
AKS-HELM / AKS-HELM2 (×8)	AKS	multiple	MEDIUM	Kubernetes cluster access
OpenAI ×3	Cognitive	various	MEDIUM	Unauthorized LLM usage

Risk Assessment Matrix



Contributing Risk Factors:

-  Tagged as **Critical Asset** in Exposure Graph
-  Active Insider Risk Incident (#49130)
-  60+ directly reachable resources with zero lateral movement required
-  SOC infrastructure within blast radius (meta-attack potential)
-  Cross-identity linkage to second compromised identity (`elviaa@zava-corp.com`)
-  Identity Protection not enrolled — no automated risk signals
-  Privileged CLI/API tooling in active use
-  35% sign-in failure rate consistent with automated tooling or credential stuffing
-  Non-corporate IP `76.182.132.142` in active sign-in history

Immediate Recommended Actions

Priority 1 — Contain (Within 2 Hours)

Action	Owner	Detail
Suspend account	IT/IdOps	Disable <code>u411@a.alpineskihouse.co</code> in Entra ID pending investigation
Revoke all active sessions	IT/IdOps	Revoke all refresh tokens: <code>Revoke-AzureADUserAllRefreshToken</code>
Block IP <code>76.182.132.142</code>	Network/CASB	Conditional Access policy or firewall block
Audit <code>midnightblizzardstorage</code>	Security Ops	Pull complete access log from storage blob audit — identify what was read/written
Rotate <code>MDTIACTOR</code> vault secrets	Security Ops	Immediate rotation of all secrets in this vault
Rotate <code>SentinelDemo</code> vault secrets	Security Ops	Prevent SIEM credential exposure

Priority 2 — Investigate (Within 24 Hours)

Action	Owner	Detail
Full <code>shellalpinesa</code> review	Security Ops	Review Cloud Shell storage for scripts, tokens, SSH keys uploaded or modified
<code>sharedcredsash</code> audit	Security Ops	Identify all secrets accessed in the past 30 days
PowerShell / CLI activity audit	Security Ops	Query <code>AzureActivity</code> and <code>AuditLogs</code> for all CLI/PowerShell operations by this identity
<code>CyberSOC_ML</code> integrity check	SOC Engineering	Verify no model modifications or data poisoning in the last 30 days
Cross-tenant zavabot investigation	Security Ops	Coordinate with Zava Corp SOC — determine how <code>u411</code> gained permissions to <code>zavabot</code> RG
Enroll in Identity Protection	IT/IdOps	Enable risk-based sign-in and user risk policies for <code>alpineskihouse.co</code> domain

Priority 3 — Remediate (Within 72 Hours)

Action	Owner	Detail
RBAC permission review	Cloud Governance	Full audit of all direct role assignments for this identity — revoke excess permissions
Key Vault access policy review	Cloud Governance	Replace broad RBAC grants with least-privilege Access Policies per vault
Insider Risk policy review	HR / Compliance	Review and escalate Incident #49130 with new evidence
Detection rule — midnightblizzard_infra	SOC	Create analytics rule alerting on any access to midnightblizzard_infra resources
Cross-tenant RBAC audit	Cloud Governance	Audit how alpineskihouse.co identities acquired permissions in zava-corp.com tenant resources

Appendix — Investigation Methodology

Step	Tool / Data Source	Finding
1. Identity resolution	SigninLogs	UPN u411@a.alpineskihouse.co = DisplayName Denise Charron , AAD ObjectId c4e33ecf-bbb2-4948-953d-d3bfe73a7d74
2. Risk status check	AADRiskyUsers	No records — Identity Protection not monitoring this account
3. Behavioral analysis	SigninLogs (30d)	800 sign-ins, 34.75% failure rate, 35 apps including CLI/API tools, non-corporate IP
4. Exposure Graph query	graph_find_blastradius("Denise Charron")	60+ directly reachable resources across Key Vaults, AKS, Storage, AI, ML
5. Graph node resolution	graph_find_nodes(User)	Graph indexes by display name; UPN/email queries return null
6. Cross-identity correlation	Blast radius comparison	zavabot RG intersection with elviaa@zava-corp.com blast radius confirmed
7. Threat intelligence enrichment	Resource group naming analysis	midnightblizzard_infra matches Midnight Blizzard (NOBELIUM/APT29) threat actor

Report generated from live Microsoft Sentinel data — March 4, 2026
Data sources: SigninLogs, AADRiskyUsers, Microsoft Sentinel Exposure Graph (graph_find_blastradius)
For questions contact your SOC analyst